

hw1q3

March 4, 2021

0.0.1 Liam Jackson

0.0.2 BE700 ML with Andy Fan

0.0.3 HW1q3

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
# %matplotlib inline
```

1. Loading Dataset

```
[2]: data_df = pd.read_csv('data.csv')
num_cols = data_df.columns.size
print(f"There are {num_cols} columns in the raw dataframe. There is one Unnamed_
↳column, the 33rd column, which contains no data.")
for ind, name in enumerate(data_df.columns):
    print(f"Column number: {ind + 1}, Column name: {name}")
```

There are 33 columns in the raw dataframe. There is one Unnamed column, the 33rd column, which contains no data.

Column number: 1, Column name: id
Column number: 2, Column name: diagnosis
Column number: 3, Column name: radius_mean
Column number: 4, Column name: texture_mean
Column number: 5, Column name: perimeter_mean
Column number: 6, Column name: area_mean
Column number: 7, Column name: smoothness_mean
Column number: 8, Column name: compactness_mean
Column number: 9, Column name: concavity_mean
Column number: 10, Column name: concave points_mean
Column number: 11, Column name: symmetry_mean
Column number: 12, Column name: fractal_dimension_mean
Column number: 13, Column name: radius_se
Column number: 14, Column name: texture_se
Column number: 15, Column name: perimeter_se
Column number: 16, Column name: area_se
Column number: 17, Column name: smoothness_se

```

Column number: 18, Column name: compactness_se
Column number: 19, Column name: concavity_se
Column number: 20, Column name: concave points_se
Column number: 21, Column name: symmetry_se
Column number: 22, Column name: fractal_dimension_se
Column number: 23, Column name: radius_worst
Column number: 24, Column name: texture_worst
Column number: 25, Column name: perimeter_worst
Column number: 26, Column name: area_worst
Column number: 27, Column name: smoothness_worst
Column number: 28, Column name: compactness_worst
Column number: 29, Column name: concavity_worst
Column number: 30, Column name: concave points_worst
Column number: 31, Column name: symmetry_worst
Column number: 32, Column name: fractal_dimension_worst
Column number: 33, Column name: Unnamed: 32

```

2. Generating a matrix scatter plot

```

[3]: mean_cols = [mean_col for mean_col in data_df.columns if 'mean' in mean_col]
mean_cols = data_df.filter(regex = 'mean').columns
print(f"Just verifying there are {len(mean_cols)} columns with 'mean' in the_
      ↪column name")

```

Just verifying there are 10 columns with 'mean' in the column name

```


[4]: means_df = data_df.filter(items = mean_cols)
means_df.columns = means_df.columns.str.replace(' ', '_')

```

```

[5]: mat_plot = pd.plotting.scatter_matrix(means_df, figsize = (15, 15))

```



3. Calculations/Statistics

```
[6]: num_ben = data_df['diagnosis'].value_counts()['B']
num_mal = data_df['diagnosis'].value_counts()['M']
print(f"There are {num_ben} Benign occurrences and {num_mal} Malignant_
↪ occurrences.")
```

There are 357 Benign occurrences and 212 Malignant occurrences.

```
[7]: stats_df = means_df.describe().loc[["mean", "std"], :]
print("Statistics for both diagnoses: ")
stats_df
```

Statistics for both diagnoses:

```
[7]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean \
mean      14.127292      19.289649      91.969033  654.889104      0.096360
std        3.524049       4.301036      24.298981  351.914129      0.014064

      compactness_mean  concavity_mean  concave_points_mean  symmetry_mean \
mean           0.104341      0.088799      0.048919      0.181162
std           0.052813      0.079720      0.038803      0.027414

      fractal_dimension_mean
mean              0.062798
std              0.007060
```

```
[8]: ben_df = data_df[data_df['diagnosis'] == 'B'].filter(items = mean_cols)
ben_stats_df = ben_df.describe().loc[["mean", "std"], :]
print("Benign Occurrences Statistics: ")
ben_stats_df
```

Benign Occurrences Statistics:

```
[8]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean \
mean      12.146524      17.914762      78.075406  462.790196      0.092478
std        1.780512       3.995125      11.807438  134.287118      0.013446

      compactness_mean  concavity_mean  concave points_mean  symmetry_mean \
mean           0.080085      0.046058      0.025717      0.174186
std           0.033750      0.043442      0.015909      0.024807

      fractal_dimension_mean
mean              0.062867
std              0.006747
```

```
[9]: mal_df = data_df[data_df['diagnosis'] == 'M'].filter(items = mean_cols)
mal_stats_df = mal_df.describe().loc[["mean", "std"], :]
print("Malignant Occurrences Statistics: ")
mal_stats_df
```

Malignant Occurrences Statistics:

```
[9]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean \
mean      17.462830      21.604906     115.365377  978.376415      0.102898
std        3.203971       3.779470      21.854653  367.937978      0.012608

      compactness_mean  concavity_mean  concave points_mean  symmetry_mean \
mean           0.145188      0.160775      0.087990      0.192909
std           0.053987      0.075019      0.034374      0.027638

      fractal_dimension_mean
```

```
mean          0.062680
std           0.007573
```

```
[10]: num_ben_rad15 = ben_df[ben_df['radius_mean'] >= 15].shape[0]
per_ben_rad15 = round(100 * (num_ben_rad15 / num_ben), 2)
print(f"{per_ben_rad15} % of Benign occurrences have a cell radius of at least_
↪15")
```

3.64 % of Benign occurrences have a cell radius of at least 15

4. Building OLS Model to predict area (y) given radius (x)

```
[11]: xy_df = data_df[['radius_mean', 'area_mean']].sort_values('radius_mean')
xy_df
```

```
[11]:      radius_mean  area_mean
101         6.981        143.5
539         7.691        170.4
538         7.729        178.8
568         7.760        181.0
46          8.196        201.9
..          ...          ...
82         25.220       1878.0
352         25.730       2010.0
180         27.220       2250.0
461         27.420       2501.0
212         28.110       2499.0
```

[569 rows x 2 columns]

```
[12]: max_poly_ord = 2
poly_ord_range = list(range(1, max_poly_ord + 1))

ols_all_models_dict = {}

for model_ord in poly_ord_range:
    model_dict = {}
    model_df = xy_df.copy()
    x = model_df['radius_mean']
    y_r = model_df['area_mean']

    X = np.zeros([len(x), model_ord + 1])

    for X_col in list(range(0, X.shape[1])):
        X[:,X_col] = x ** X_col
    Xt = np.transpose(X)
    XtX = np.matmul(Xt, X)
    XtXinv = np.linalg.inv(XtX)
```

```

XtXinvXt = np.matmul(XtXinv, Xt)
beta = np.matmul(XtXinvXt, y_r)
beta_flip = np.flipud(beta)

y_m = np.polyval(beta_flip, x)
res = y_r.subtract(y_m)
res_sq = abs(res) ** 2

model_df['y_m'] = y_m
model_df['res'] = res
model_df['res_sq'] = res_sq

res_sq_sum = sum(res_sq)
MSE = res_sq_sum / len(x)

model_dict['beta_flip'] = beta_flip
model_dict['model_df'] = model_df
model_dict['res_sq_sum'] = res_sq_sum
model_dict['MSE'] = MSE

model_key = "p = " + str(model_ord)
ols_all_models_dict[model_key] = model_dict

```

```
[13]: ols_all_models_dict
```

```

[13]: {'p = 1': {'beta_flip': array([ 98.59821922, -738.0367042 ]),
  'model_df':      radius_mean  area_mean          y_m          res
res_sq
101          6.981          143.5   -49.722536   193.222536   37334.948362
539          7.691          170.4    20.282200   150.117800   22535.353941
538          7.729          178.8    24.028932   154.771068   23954.083453
568          7.760          181.0    27.085477   153.914523   23689.680417
46           8.196          201.9    70.074300   131.825700   17378.015051
..           ...           ...           ...           ...
82           25.220         1878.0   1748.610384   129.389616   16741.672622
352          25.730         2010.0   1798.895476   211.104524   44565.119965
180          27.220         2250.0   1945.806823   304.193177   92533.489030
461          27.420         2501.0   1965.526467   535.473533   286731.904882
212          28.110         2499.0   2033.559238   465.440762   216635.102985

[569 rows x 5 columns],
'res_sq_sum': 1767428.9562542248,
'MSE': 3106.2020320812385},
'p = 2': {'beta_flip': array([ 3.10992516,  0.43684601, -10.5164038 ]),
  'model_df':      radius_mean  area_mean          y_m          res          res_sq
101          6.981          143.5   144.093434   -0.593434    0.352164

```

539	7.691	170.4	176.800058	-6.400058	40.960745
538	7.729	178.8	178.638950	0.161050	0.025937
568	7.760	181.0	180.145751	0.854249	0.729742
46	8.196	201.9	201.971393	-0.071393	0.005097
..
82	25.220	1878.0	1978.563778	-100.563778	10113.073432
352	25.730	2010.0	2059.596420	-49.596420	2459.804862
180	27.220	2250.0	2305.606421	-55.606421	3092.074085
461	27.420	2501.0	2339.679053	161.320947	26024.448049
212	28.110	2499.0	2459.139436	39.860564	1588.864558

```
[569 rows x 5 columns],
'res_sq_sum': 123097.70230710595,
'MSE': 216.34042584728638}}
```

```
[14]: print(f"The Linear (p = 1) model coefficients are: {ols_all_models_dict['p = 1']['beta_flip']}")
print("The Linear (p = 1) model residuals are: ")
print(ols_all_models_dict['p = 1']['model_df']['res'])
```

The Linear (p = 1) model coefficients are: [98.59821922 -738.0367042]

The Linear (p = 1) model residuals are:

101	193.222536
539	150.117800
538	154.771068
568	153.914523
46	131.825700

...	
82	129.389616
352	211.104524
180	304.193177
461	535.473533
212	465.440762

Name: res, Length: 569, dtype: float64

```
[15]: print(f"The Quadratic (p = 2) model coefficients are: {ols_all_models_dict['p = 2']['beta_flip']}")
print(f"The Quadratic (p = 2) model residuals are: ")
print(ols_all_models_dict['p = 2']['model_df']['res'])
```

The Quadratic (p = 2) model coefficients are: [3.10992516 0.43684601 -10.5164038]

The Quadratic (p = 2) model residuals are:

101	-0.593434
539	-6.400058
538	0.161050
568	0.854249

```

46      -0.071393
      ...
82     -100.563778
352     -49.596420
180     -55.606421
461     161.320947
212      39.860564
Name: res, Length: 569, dtype: float64

```

5. Plotting Data vs. Polynomial Models

```

[16]: p1_coeffs = ols_all_models_dict['p = 1']['beta_flip']
      p2_coeffs = ols_all_models_dict['p = 2']['beta_flip']

      p1_data_df = ols_all_models_dict['p = 1']['model_df']
      p2_data_df = ols_all_models_dict['p = 2']['model_df']

      x_r = p1_data_df['radius_mean']
      y_r = p1_data_df['area_mean']
      p1_y_m = p1_data_df['y_m']
      p2_y_m = p2_data_df['y_m']

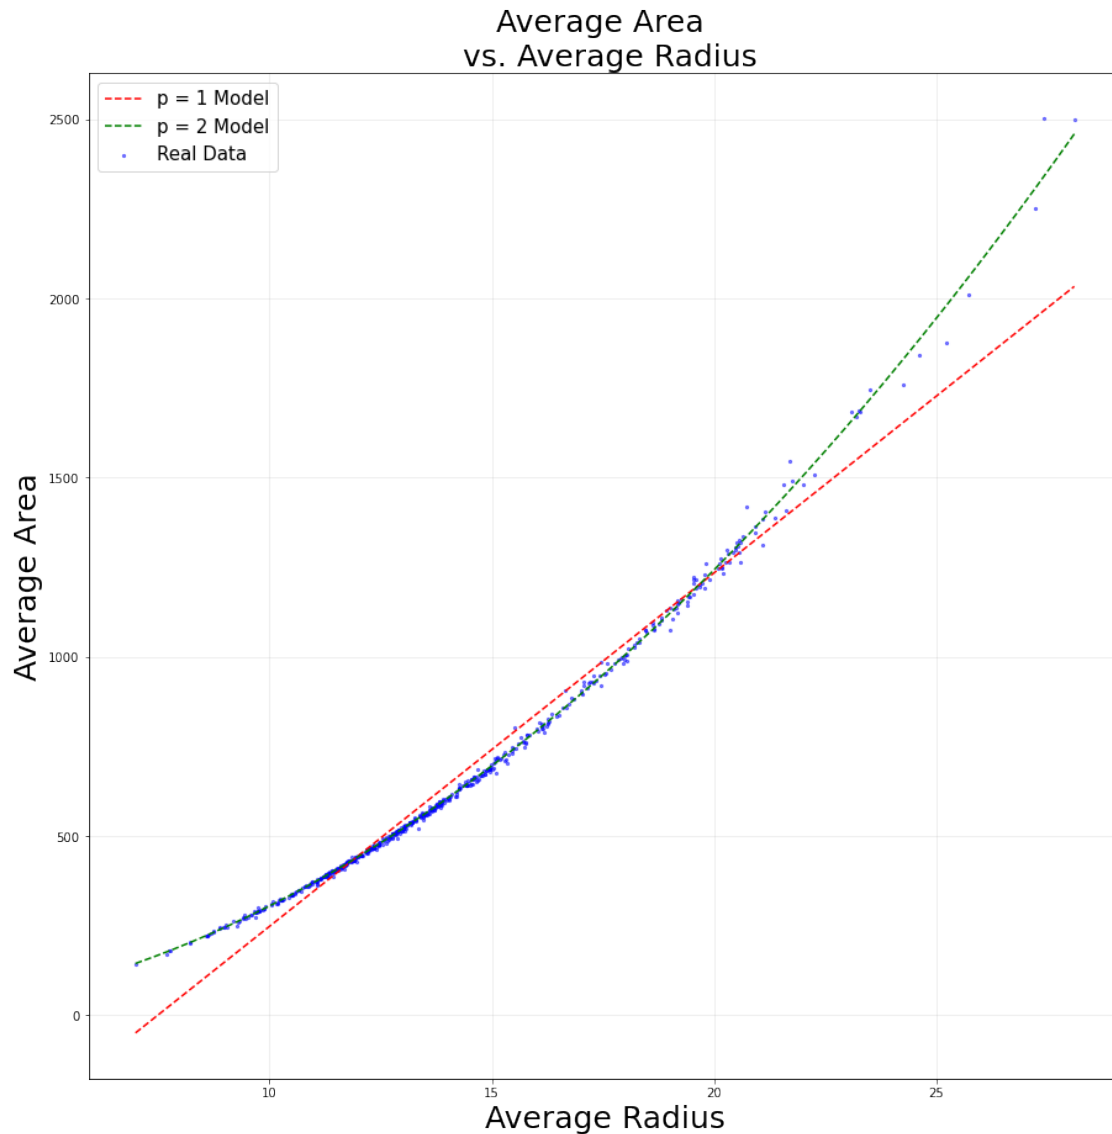
      x_m = np.linspace(min(x_r), max(x_r))
      p1_y_m_linspace = np.polyval(p1_coeffs, x_m)
      p2_y_m_linspace = np.polyval(p2_coeffs, x_m)

[17]: plt.figure(figsize = (15, 15))

      plt.scatter(x_r, y_r, s=5, c='b', alpha=0.5, label = 'Real Data')
      plt.plot(x_m, p1_y_m_linspace, 'r--', label = 'p = 1 Model')
      plt.plot(x_m, p2_y_m_linspace, 'g--', label = 'p = 2 Model')
      plt.grid(alpha = .25)

      plt.title('Average Area \n vs. Average Radius', fontsize = 25)
      plt.legend(fontsize = 15)
      plt.xlabel('Average Radius', fontsize = 25)
      plt.ylabel('Average Area', fontsize = 25)
      plt.show()

```

```
[18]: plt.figure(figsize = (15, 15))

plt.subplot(2, 1, 1)

for i in range(len(x_r)):
    p1_res_x = (x_r[i], x_r[i])
    p1_res_y = (y_r[i], p1_y_m[i])
    plt.plot(p1_res_x, p1_res_y, color = 'orange', linewidth = 2, alpha = 0.
↪5)

plt.scatter(x_r, y_r, s=10, c='b', alpha=0.5, label = 'Real Data')
plt.plot(x_m, p1_y_m_linspace, 'r--', label = 'p = 1 Model')
```

```

plt.grid(alpha = .25)
plt.title('Residuals in Polynomial Models for Average Area \n vs. Average_
↪Radius', fontsize = 25)
plt.legend(fontsize = 15, loc = 'lower right')
plt.ylabel('Average Area', fontsize = 25)
plt.xlim([20, 30])
plt.ylim([1000, 2750])

plt.subplot(2, 1, 2)

for i in range(len(x_r)):
    p2_res_x = (x_r[i], x_r[i])
    p2_res_y = (y_r[i], p2_y_m[i])
    plt.plot(p2_res_x, p2_res_y, color = 'orange', linewidth = 2, alpha = 0.
↪5)

plt.scatter(x_r, y_r, s=10, c='b', alpha=0.5, label = 'Real Data')
plt.plot(x_m, p2_y_m_linspace, 'g--', label = 'p = 2 Model')

plt.grid(alpha = .25)
plt.legend(fontsize = 15, loc = 'lower right')
plt.xlabel('Average Radius', fontsize = 25)
plt.ylabel('Average Area', fontsize = 25)
plt.xlim([20, 30])
plt.ylim([1000, 2750])

plt.show()

```

Residuals in Polynomial Models for Average Area
vs. Average Radius

